

■ Aprendizaje por refuerzo con Open AI Gym

Ana Laura Diedrichs @anadiedrichs

Link a esta presentación:

<https://tinyurl.com/pydatarl>



Hola!

- Ph.D. candidate at UTN FRM
- Software engineer -
- IoT ML consultant
- Teacher assistant
- ana.diedrichs@frm.utn.edu.ar
- Web <https://anadiedrichs.github.io/>
- Twitter @anadiedrichs




Motivación



Supervised learning

Aprendizaje supervisado



https://commons.wikimedia.org/wiki/File:Supervised_machine_learning_in_a_nutshell.svg

CLASSICAL MACHINE LEARNING

Data is pre-categorized
or numerical

SUPERVISED

Predict
a category

CLASSIFICATION

«Divide the socks by color»



Predict
a number

REGRESSION

«Divide the ties by length»




Data is not labeled
in any way

UNSUPERVISED

Divide
by similarity

CLUSTERING

«Split up similar clothing
into stacks»



Identify sequences

Find hidden
dependencies

ASSOCIATION

«Find what clothes I often
wear together»




DIMENSION REDUCTION (generalization)

«Make the best outfits from the given clothes»



Reinforcement learning (RL)

Aprendizaje por refuerzo



https://commons.wikimedia.org/wiki/File:RL_agent.png

Reinforcement learning (RL)

Aprendizaje por refuerzo



https://en.m.wikipedia.org/wiki/File:Learning_to_walk.png

Reinforcement learning (RL)

Aprendizaje por refuerzo




Figure 3.1: The agent–environment interaction in a Markov decision process.


Source: [Reinforcement Learning: An Introduction](#)

RL applications

Aplicaciones



(a) OBELIX robot




(b) Zebra Zero robot



(c) Autonomous helicopter



(d) Sarcos humanoid
DB




J. Kober, J. A. D. Bagnell, J. Peters. Reinforcement Learning in Robotics: A survey. Int. J. Robot. Res. Jul. 2013.
https://www.ias.informatik.tu-darmstadt.de/uploads/Publications/Kober_IJRR_2013.pdf

RL applications

Aplicaciones






RL vs linear model vs Human. Click [here](#) for the source
<https://web.stanford.edu/class/psych209/Readings/MnihEtAlHassibis15NatureControlDeepRL.pdf>

RL applications


Alpha Go zero



<https://deepmind.com/blog/article/alphago-zero-starting-scratch>

RL applications

Aplicaciones



<https://deepmind.com/blog/article/alphago-zero-starting-scratch>

RL applications


Alpha Go zero



<https://deepmind.com/blog/article/alphago-zero-starting-scratch>

RL applications

Alpha Go zero



<https://deepmind.com/blog/article/alphago-zero-starting-scratch>

RL applications

Aplicaciones

Z. Zhou, X. Li, and R. N. Zare. Optimizing Chemical Reactions with Deep Reinforcement Learning. ACS Central Science 3, 2017.

<https://github.com/lightingghost/chemopt>



RL applications

Aplicaciones

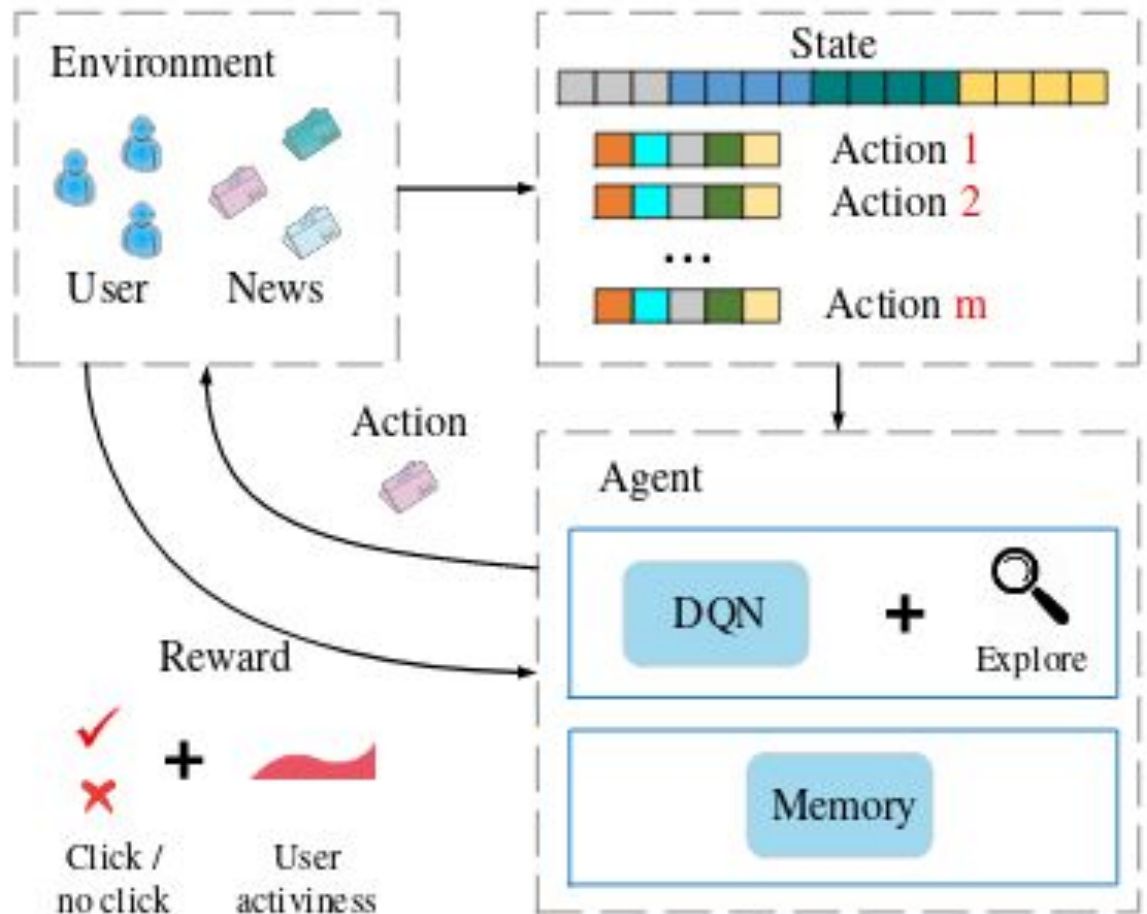


Figure 2: Deep Reinforcement Recommendation System

- G. Zheng, F. Zhang, Z. Zheng, Y. Xiang, Ni. J. Yuan, X. Xie, and Z. Li. DRN: A Deep Reinforcement Learning Framework for News Recommendation. 2018.



RL applications

Aplicaciones

- Polydoros, Athanasios S., and Lazaros Nalpantidis. "Survey of model-based reinforcement learning: Applications on robotics." *Journal of Intelligent & Robotic Systems* 86.2 (2017): 153-173.
- H.Mao, Alizadeh, M. Alizadeh, Menache, I.Menache, and S.Kandula. Resource Management With deep Reinforcement Learning. In *ACM Workshop on Hot Topics in Networks*, 2016.
- X. Bu, J. Rao, C. Z. Xu. A reinforcement learning approach to online web systems auto-configuration. In *Distributed Computing Systems*, 2009. *ICDCS'09.29th IEEE International Conference on*. IEEE , 2019.
- J. Jin, C.Song, H. Li, K. Gai, J.Wang amd W. Zhang. Real-Time Bidding with Multi-Agent Reinforcement Learning in Display Advertising. *arXiv preprint arXiv:1802.09756*, 2018.

Reproducibility issues

Problemas de reproducibilidad

- How to compare?
- ¿Cómo comparamos? Criterios



Open AI Gym




Gym

Gym is a toolkit for developing and comparing reinforcement learning algorithms. It supports teaching agents everything from walking to playing games like Pong or Pinball.

[View documentation >](#)

[View on GitHub >](#)



RandomAgent on CartPole-v1

<https://gym.openai.com/>

<https://openai.com>

Box2D

Continuous control tasks in the Box2D simulator.



BipedalWalker-v2
Train a bipedal robot to walk.



BipedalWalkerHardcore-v2
Train a bipedal robot to walk over rough terrain.



CarRacing-v0
Race a car around a track.



LunarLander-v2
Navigate a lander to its landing pad.



LunarLanderContinuous-v2
Navigate a lander to its landing pad.

Robotics

Simulated **goal-based tasks** for the Fetch and ShadowHand robots.



FetchPickAndPlace-v1
Lift a block into the air.



FetchPush-v1
Push a block to a goal position.




FetchReach-v1
Move Fetch to a goal position.



FetchSlide-v1
Slide a puck to a goal position.



HandManipulateBlock-v0
Orient a block using a robot hand.




HandManipulateEgg-v0
Orient an egg using a robot hand.

Classic control


Control theory problems from the classic RL literature.




Acrobot-v1
Swing up a two-link robot.



CartPole-v1
Balance a pole on a cart.



MountainCar-v0
Drive up a big hill.



MountainCarContinuous-v0
Drive up a big hill with continuous control.



Pendulum-v0
Swing up a pendulum.




AirRaid-ram-v0
Maximize score in the game
AirRaid, with RAM as input



AirRaid-v0
Maximize score in the game
AirRaid, with screen images
as input



Alien-ram-v0
Maximize score in the game
Alien, with RAM as input



Open AI Gym



Open source interface to reinforcement learning tasks.

The `gym` library provides an easy-to-use suite of reinforcement learning tasks.

```
import gym
env = gym.make("CartPole-v1")
observation = env.reset()
for _ in range(1000):
    env.render()
    action = env.action_space.sample() # your agent here (this takes random actions)
    observation, reward, done, info = env.step(action)

    if done:
        observation = env.reset()
env.close()
```

<https://gym.openai.com/>

Demo: pacman aleatorio

- Source:
<https://star-ai.github.io/Rendering-OpenAi-Gym-in-Colaboratory/>
- <https://colab.research.google.com/drive/19cNaxTcwQ-UjAHmjoTDQVtbSSROGG616>

Demo: taxi environment

- Taxi - Q-learning
- https://colab.research.google.com/drive/1mNn9Ij5ue0W66gjn4_fnK6BIPLKHg8wt
- Frozen lake iteración de valores
https://colab.research.google.com/drive/12LVBVI_C4TrNDV2bMgW-fnuWPeoCrure

Q-learning

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

Initialize Q table



Choose an action a



Perform action



Measure reward




Update Q



At the end of the training




Good Q*table




Deep RL

Action selection step



Training step



¿Aplicar RL en mi proyecto?

- May I use RL in my project?


Understand, study
your
problem!



¿Aplicar RL en mi proyecto?



RL requirements



- Que debo tener en cuenta antes de aplicar RL
 - MDP definition: definir acciones, estados, recompensa, etc
 - Simulated environment: simular el entorno del agente.

Recursos para aprender

- Implementation of Reinforcement Learning Algorithms. <https://github.com/dennybritz/reinforcement-learning>
- A curated list of resources dedicated to reinforcement learning. <https://github.com/aikorea/awesome-rl>
- Experimentos con RL: SLM lab
<https://kengz.gitbooks.io/slm-lab/content/>
- Ejemplo en robótica con Arduino
<https://www.michaelang.com/project/sisyphus>
- https://medium.com/@robertmoni_66330/reinforcement-learning-algorithms-an-intuitive-overview-of-existing-algorithms-c2095902867a
- Notebook RL Qlearning ejemplo curso
<https://colab.research.google.com/drive/1fl5bZtO4xhJNR34Zayyrpqq2vkeGug6y>

¡Muchas gracias PyData!

- Give me feedback :-) [en este formulario](#)
- ¿Preguntas?
- Ana Laura Diedrichs @anadiedrichs
- Ana.diedrichs@frm.utn.edu.ar

